

Jaeger trace 文档

打点代码分支：

<https://github.com/pyxyzc/sglang/tree/jeager>

安装依赖服务

安装 启动 Jaeger (UI)

代码块

```
1 wget https://github.com/jaegertracing/jaeger/releases/download/v1.76.0/jaeger-  
1.76.0-linux-amd64.tar.gz  
2  
3 tar -xzvf jaeger-1.76.0-linux-amd64.tar.gz  
4  
5 cd jaeger-1.76.0-linux-amd64  
6  
7 # 使用绝对路径进行软连接保证全局可用, 修改为jaeger-all-in-one的绝对路径  
8 ln -s /sgl-workspace/jaeger/jaeger-1.76.0-linux-amd64/jaeger-all-in-one \  
9 /usr/local/bin/jaeger-all-in-one  
10  
11 jaeger-all-in-one --collector.otlp.enabled=true --collector.otlp.grpc.host-  
port=0.0.0.0:4327 --collector.otlp.http.host-port=0.0.0.0:4328
```

报错 too many open files，具体报错信息：

```
{"level":"panic","ts":1765965126.7292247,"caller":"app/static_handler  
.go:46","msg":"Could not create static assets handler","error":"too  
many open files","stacktrace": ...}
```

```
panic: Could not create static assets handler
```

解决方法如下：

编辑 /etc/security/limits.conf

代码块

```
1 jaeger_user soft nofile 65536  
2 jaeger_user hard nofile 65536
```

安装启动 otel-collector (数据中转)

代码块

```
1 wget https://github.com/open-telemetry/opentelemetry-collector-
2   releases/releases/download/v0.114.0/otelcol-contrib_0.114.0_linux_amd64.tar.gz
3 tar -xzvf otelcol-contrib_0.114.0_*.tar.gz
4
5 ln -s otelcol-contrib /usr/local/bin/otelcol-contrib
6
7
8 # 启动前先创建配置文件
9 otelcol-contrib --config otel-collector.yaml
```

otel-collector.yaml

```
1 receivers:
2   otlp:
3     protocols:
4       grpc:
5         endpoint: 0.0.0.0:4317
6       http:
7         endpoint: 0.0.0.0:4318
8
9 exporters:
10  otlp:
11    endpoint: "localhost:4327"      # Jaeger OTLP gRPC 地址
12    tls:
13      insecure: true            # 本地环境常用
14
15 service:
16  pipelines:
17    traces:
18      receivers: [otlp]
19      exporters: [otlp]
20
```

Sglang 启动参数

代码块

```
1 #!/bin/bash
```

```
2
3 # 模型路径 (改成你本地的, 比如 /data/models/Llama-2-7b-hf)
4 MODEL_PATH=/home/models/QwQ-32B
5
6 # 端口号 (可以改, 比如 30000)
7 PORT=30000
8 # 卡数量
9 TP=2
10
11 # 只使用
12 export CUDA_VISIBLE_DEVICES=0,1
13 # 启动命令
14 python python/sglang/launch_server.py \
15     --model-path $MODEL_PATH \
16     --tp $TP \
17     --port $PORT \
18     --enable-trace \
19     --otlp-traces-endpoint 127.0.0.1:4317 \
20     # --enable-metrics
```

同时启停脚本

需要将otel-collector启动配置文件放到sglang文件夹内，直接运行即可。

run_server_jaeger.sh

```
1 #!/usr/bin/env bash
2 set -e
3
4 #-----
5 # 端口检查函数
6 #-----
7 check_and_kill_port() {
8     local port=$1
9     local pid
10
11     pid=$(lsof -t -i:"$port" || true)
12     if [[ -n "$pid" ]]; then
13         echo "Port $port is in use by PID $pid. Killing..."
14         kill -9 "$pid" || true
15         echo "Port $port cleared."
16     else
17         echo "Port $port is free."
18     fi
19 }
```


```
20
21 #-----
22 # 启动前的环境变量
23 #-----
24 export SGLANG_HICACHE_FILE_BACKEND_STORAGE_DIR=/sgl-workspace/hicachefile
25
26 MODEL_PATH=/home/models/QwQ-32B
27 PORT=30000
28 TP=2
29
30 export CUDA_VISIBLE_DEVICES=0,1
31
32 #-----
33 # 启动前检查端口
34 #-----
35 echo "Checking required ports..."
36 check_and_kill_port 4317    # otlp grpc
37 check_and_kill_port 4327    # jaeger grpc
38 check_and_kill_port 4328    # jaeger http
39 check_and_kill_port $PORT   # sglang server
40
41 #-----
42 # 捕获退出信号，统一杀死所有子进程
43 #-----
44 pids=()
45
46 cleanup() {
47     echo "Received stop signal. Killing all subprocesses..."
48     for pid in "${pids[@]}"; do
49         if kill -0 "$pid" 2>/dev/null; then
50             kill -9 "$pid" 2>/dev/null || true
51         fi
52     done
53     exit 1
54 }
55
56 trap cleanup SIGINT SIGTERM EXIT
57
58 #-----
59 # 启动第 1 个: Jaeger
60 #-----
61 jaeger-all-in-one \
62     --collector.otlp.enabled=true \
63     --collector.otlp.grpc.host-port=0.0.0.0:4327 \
64     --collector.otlp.http.host-port=0.0.0.0:4328 &
65 pids+=($!)
66 echo "Started Jaeger, PID=${pids[-1]}"
```

```
67
68 #-----
69 # 启动第 2 个: otelcol-contrib
70 #-----
71 otelcol-contrib --config otel-collector.yaml &
72 pids+=($!)
73 echo "Started otelcol-contrib, PID=${pids[-1]}"
74
75 #-----
76 # 启动第 3 个: sclang server
77 #-----
78 python python/sclang/launch_server.py \
79     --model-path $MODEL_PATH \
80     --page-size 128 \
81     --tp $TP \
82     --port $PORT \
83     --enable-hierarchical-cache \
84     --hicache-write-policy write_through \
85     --hicache-storage-backend file \
86     --hicache-storage-prefetch-policy wait_complete \
87     --enable-trace \
88     --otlp-traces-endpoint 127.0.0.1:4317 &
89 pids+=($!)
90 echo "Started sclang server, PID=${pids[-1]}"
91
92 #-----
93 # 等待任意子进程退出
94 #-----
95 echo "All processes started. Waiting for any to exit..."
96 wait -n
97
98 echo "One process exited. Killing all..."
99 cleanup
100
```

结果查看


VScode 转发端口

16686 端口




查看打点数据!


scheduler调度数据



prefetch线程耗时




lookup线程耗时



dump线程耗时

对dump线程进行绑定后，可以记录dump耗时



添加打点教程

https://docs.sglang.io/references/production_request_trace.html

注意：

1. 和req相关的打点直接增加 `trace_slice_start`、`trace_slice_end` 即可